



Interprocess Communication (IPC)

- IPC methods:
 - File locking, pipes, FIFOs, message queues(system V),semaphores, shared memory.
- On a single system



Interprocess Communication (IPC)

- On different system



File and Record Locking

- Consider a line printer.
- (lock it)
- it reads the sequence number file
- it uses the number
- it increments the number and writes it back
- (unlock it)

Program

```

#define SEQFILE "seqno"          /* filename */
#define MAXBUFF  100

main()
{
    int    fd, i, n, pid, seqno;
    char    buff[MAXBUFF + 1];

    pid = getpid();
    if ( (fd = open(SEQFILE, 2)) < 0)
        err_sys("can't open %s", SEQFILE);

    for (i = 0; i < 20; i++) {
        my_lock(fd);                /* lock the file */

        lseek(fd, 0L, 0);            /* rewind before read */
        if ( (n = read(fd, buff, MAXBUFF)) <= 0)
            err_sys("read error");
        buff[n] = '\0';              /* null terminate for sscanf */

        if ( (n = sscanf(buff, "%d\n", &seqno)) != 1)
            err_sys("sscanf error");
        printf("pid = %d, seq# = %d\n", pid, seqno);

        seqno++;                    /* increment the sequence number */

        sprintf(buff, "%03d\n", seqno);
        n = strlen(buff);
        lseek(fd, 0L, 0);            /* rewind before write */
        if (write(fd, buff, n) != n)
            err_sys("write error");

        my_unlock(fd);              /* unlock the file */
    }
}

```

(lock it)

it reads the sequence

it increments the num

writes it back

(unlock it)

No Locking

```
/*
 * Locking routines that do nothing.
 */

my_lock(fd)
int    fd;
{
    return;
}

my_unlock(fd)
int    fd;
{
    return;
}
```

Results for No Locking

```
pid = 186, seq# = 1
```

(1)

```
pid = 187, seq# = 1
```

```
pid = 187, seq# = 2
```

```
pid = 187, seq# = 3
```

```
pid = 187, seq# = 4
```

```
pid = 187, seq# = 5
```

```
pid = 187, seq# = 6
```

```
pid = 187, seq# = 7
```

```
pid = 187, seq# = 8
```

```
pid = 187, seq# = 9
```

```
pid = 187, seq# = 10
```

(2)

```
pid = 186, seq# = 2
```

```
pid = 186, seq# = 3
```

```
pid = 186, seq# = 4
```

```
pid = 186, seq# = 5
```

(3)

```
pid = 187, seq# = 5
```

(4)

```
pid = 186, seq# = 6
```

```
pid = 186, seq# = 7
```

```
pid = 186, seq# = 8
```

```
pid = 186, seq# = 9
```

```
pid = 186, seq# = 10
```

```
pid = 186, seq# = 11
```

(5)

```
pid = 187, seq# = 6
```

```
pid = 187, seq# = 7
```

```
pid = 187, seq# = 8
```

```
pid = 187, seq# = 9
```

```
pid = 187, seq# = 10
```

(6)

```
pid = 186, seq# = 12
```

```
pid = 186, seq# = 13
```

```
pid = 186, seq# = 14
```

```
pid = 186, seq# = 15
```

```
pid = 186, seq# = 16
```

```
pid = 186, seq# = 17
```

```
pid = 186, seq# = 18
```

```
pid = 186, seq# = 19
```

```
pid = 186, seq# = 20
```

(7)

```
pid = 187, seq# = 11
```

```
pid = 187, seq# = 12
```

```
pid = 187, seq# = 13
```

```
pid = 187, seq# = 14
```

BSD Locking

```
/*
 * Locking routines for 4.3BSD.
 */

#include <sys/file.h>

my_lock(fd)
int fd;
{
    if (flock(fd, LOCK_EX) == -1)
        err_sys("can't LOCK_EX");
}

my_unlock(fd)
int fd;
{
    if (flock(fd, LOCK_UN) == -1)
        err_sys("can't LOCK_UN");
}
```

Results for BSD Locking

pid = 308, seq# = 1

pid = 307, seq# = 2

pid = 307, seq# = 3

pid = 307, seq# = 4

pid = 307, seq# = 5

pid = 307, seq# = 6

pid = 307, seq# = 7

pid = 307, seq# = 8

pid = 307, seq# = 9

pid = 307, seq# = 10

pid = 307, seq# = 11

pid = 307, seq# = 12

pid = 307, seq# = 13

pid = 307, seq# = 14

pid = 307, seq# = 15

pid = 308, seq# = 16

pid = 308, seq# = 17

pid = 308, seq# = 18

pid = 308, seq# = 19

pid = 308, seq# = 20

pid = 308, seq# = 21

pid = 308, seq# = 22

pid = 308, seq# = 23

pid = 308, seq# = 24

pid = 308, seq# = 25

pid = 308, seq# = 26

pid = 308, seq# = 27

pid = 307, seq# = 28

pid = 307, seq# = 29

pid = 307, seq# = 30

pid = 308, seq# = 31

pid = 308, seq# = 32

pid = 308, seq# = 33

pid = 308, seq# = 34

pid = 307, seq# = 35

pid = 307, seq# = 36

pid = 307, seq# = 37

pid = 308, seq# = 38

pid = 308, seq# = 39

pid = 308, seq# = 40

Pipes

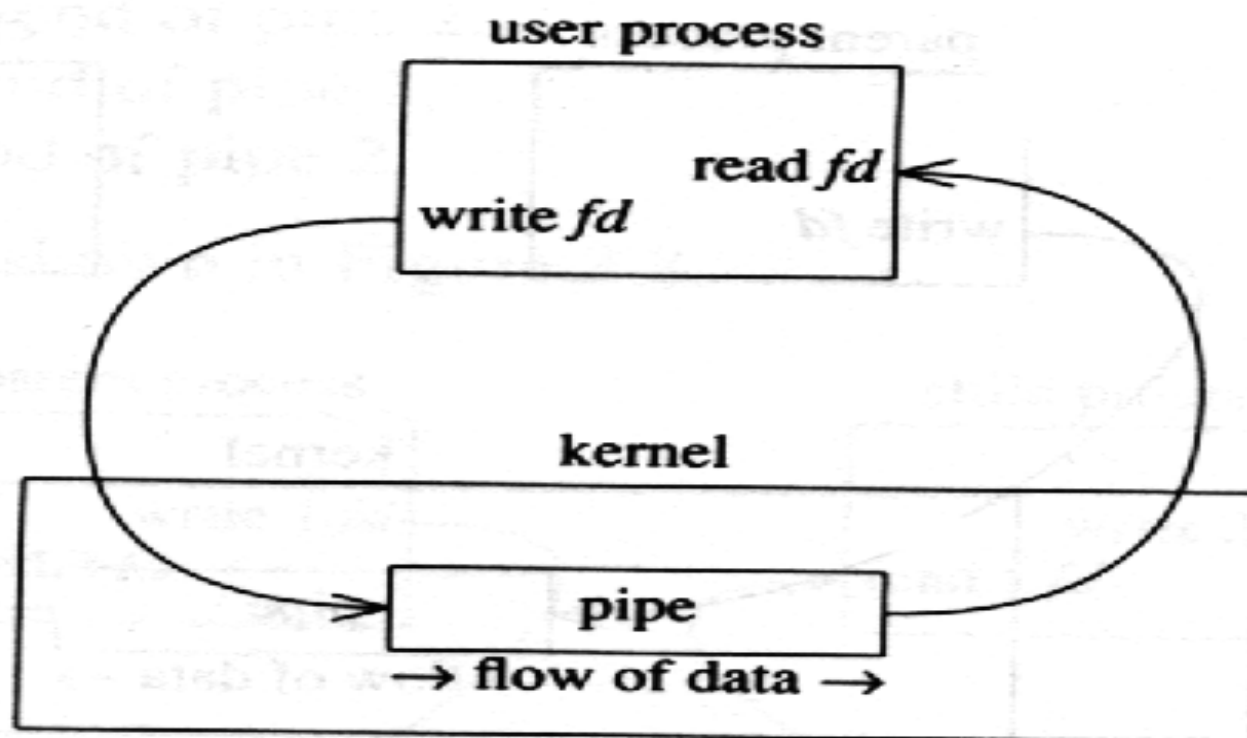


Figure 3.4 Pipe in a single process.

```
main()
{
    int    pipefd[2], n;
    char    buff[100];

    if (pipe(pipefd) < 0)
        err_sys("pipe error");

    printf("read fd = %d, write fd = %d\n", pipefd[0], pipefd[1]);
    if (write(pipefd[1], "hello world\n", 12) != 12)
        err_sys("write error");

    if ( (n = read(pipefd[0], buff, sizeof(buff))) <= 0)
        err_sys("read error");

    write(1, buff, n);        /* fd 1 = stdout */

    exit(0);
}
```

The output of this program is

```
hello world
read fd = 3, write fd = 4
```

Buffering for printf

- Save data into buffer.
- Output to device when
 - The buffer is full.
 - `fflush()` is called.
 - `exit()` is called.
 - For console output,
 - ▶ An linefeed is output.
 - ▶ The program attempts to read from the terminal.
- Good guidelines for socket programming.
 - Do not mix `printf()` and `write()`.
 - Do not use `printf()` for heavy interaction.
 - ▶ If you really want, add “flush”.
 - ▶ It is ok for CGI-like operations.

Process-to-Process Piping

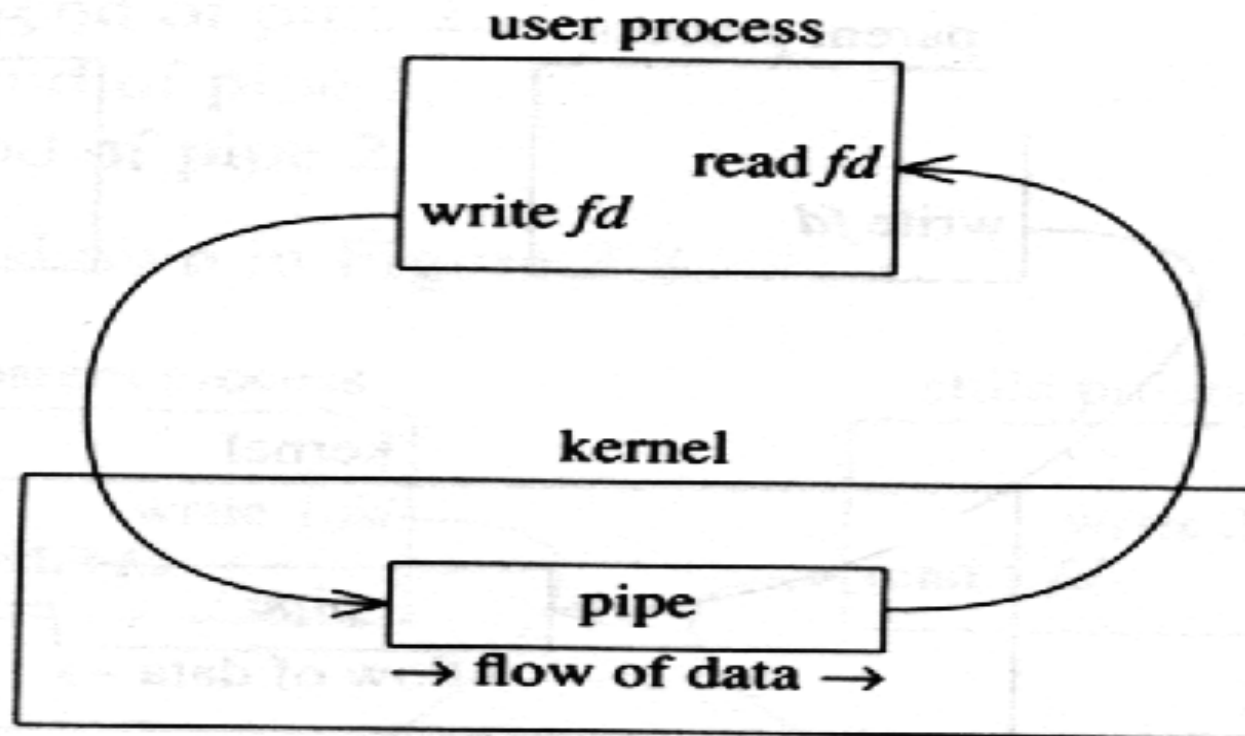
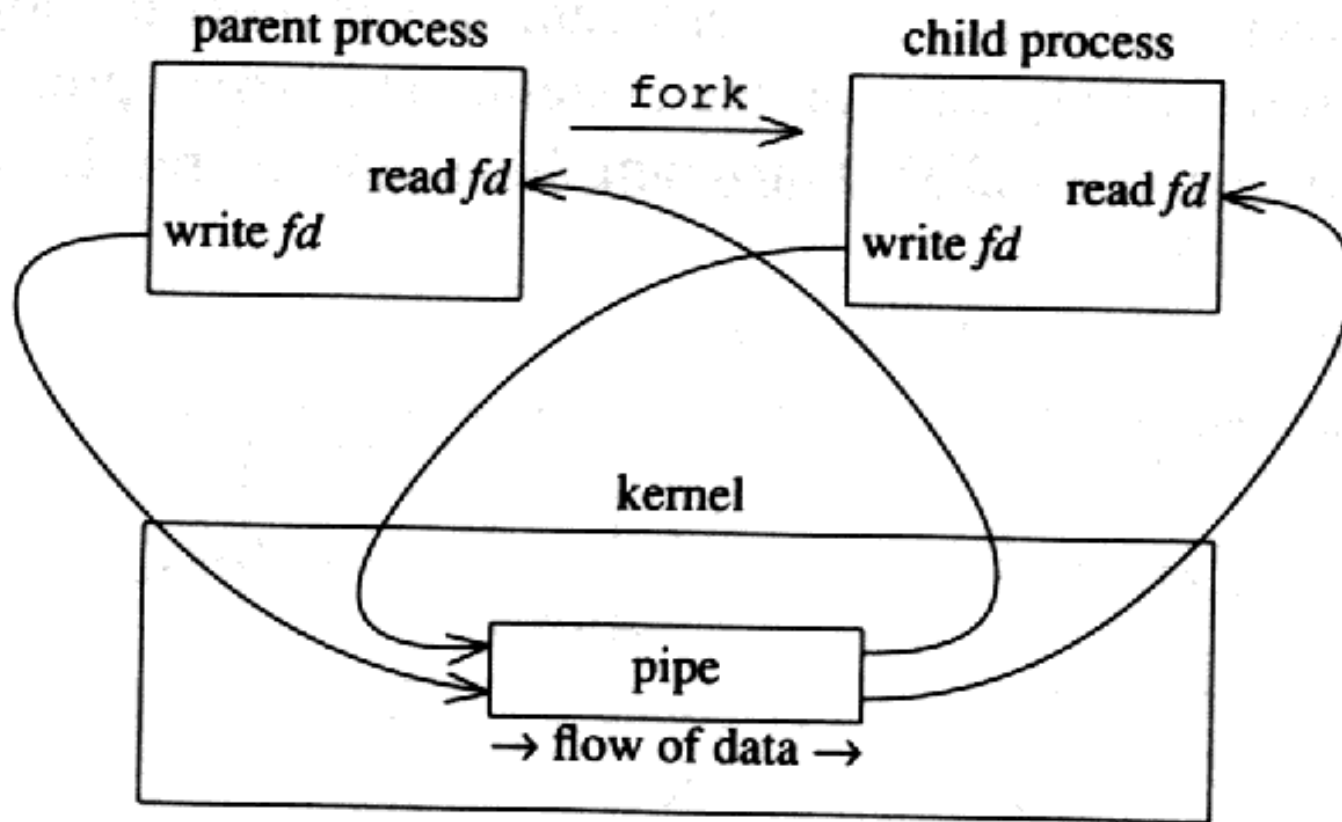
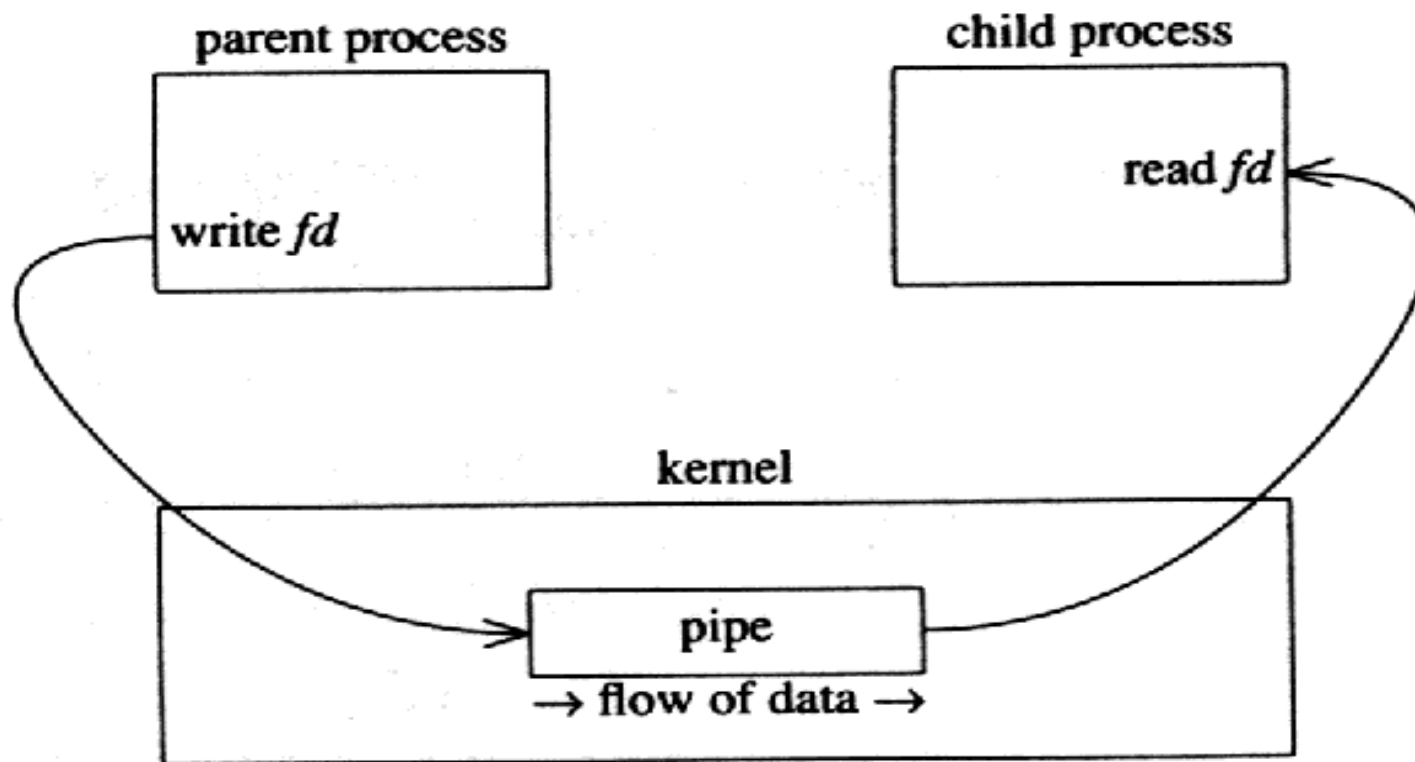
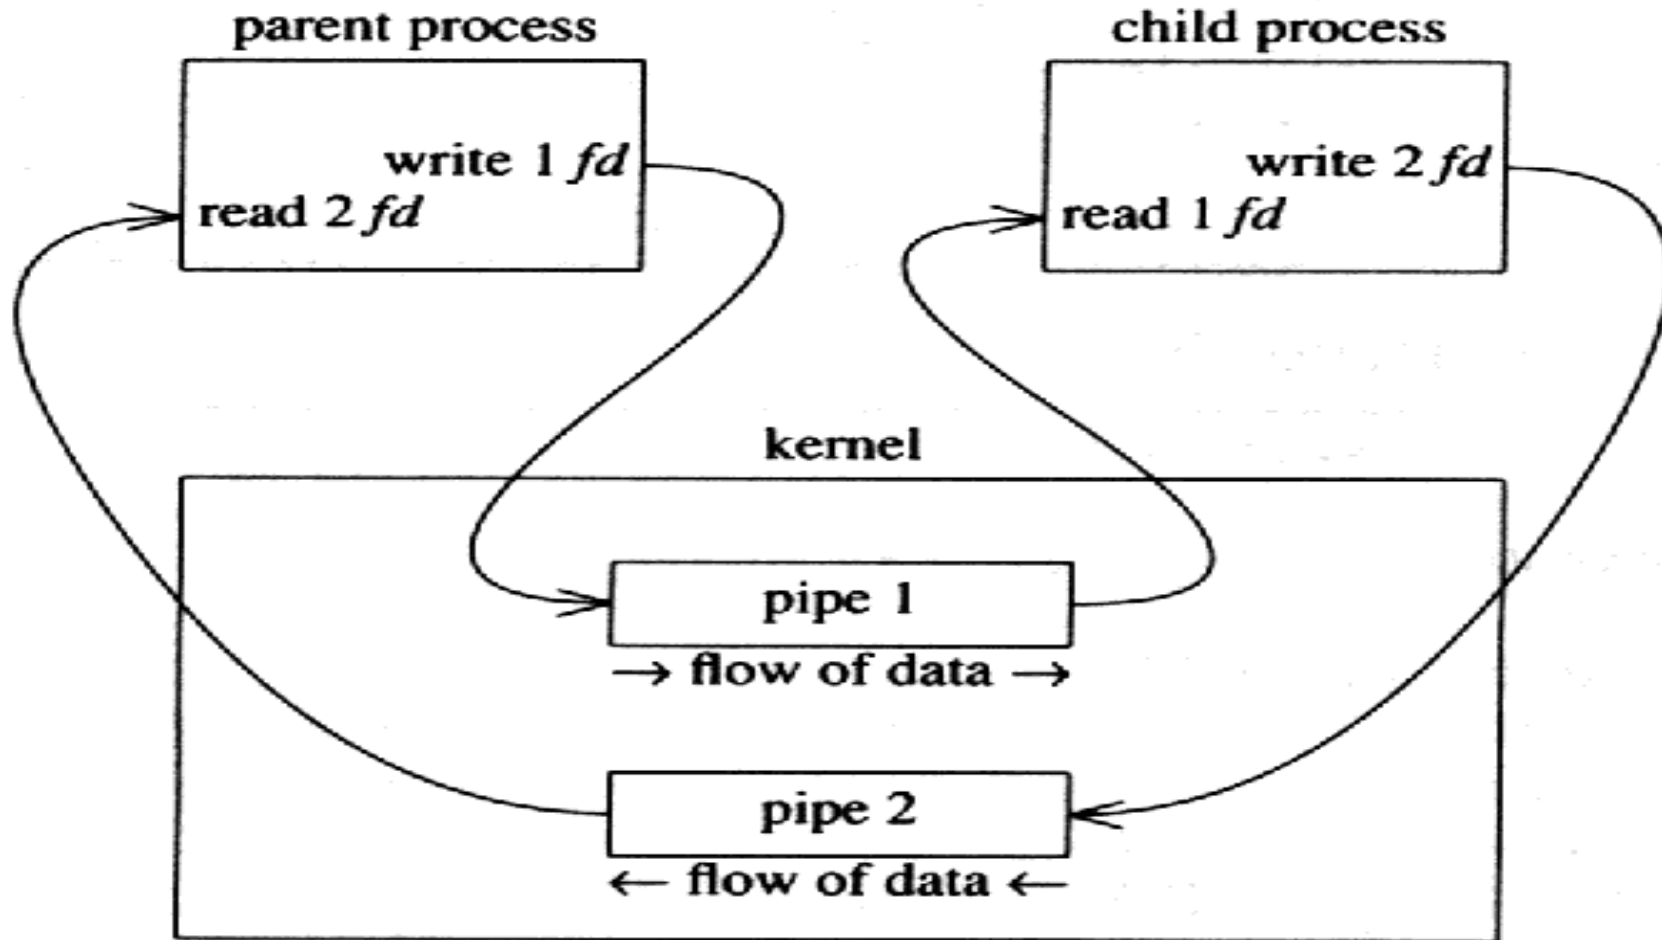


Figure 3.4 Pipe in a single process.

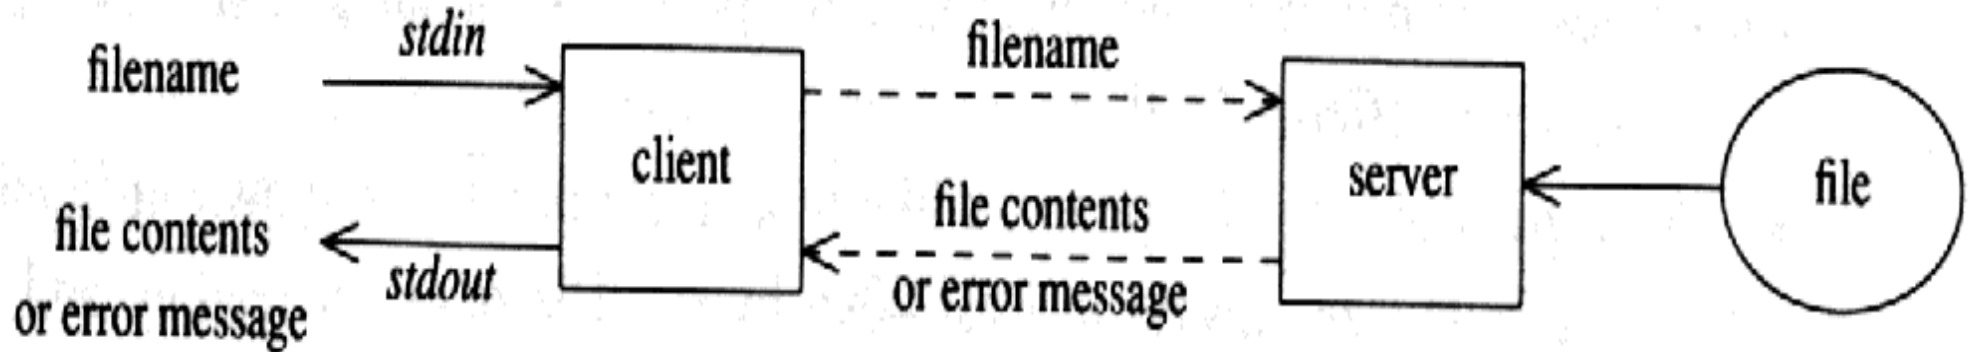




Two-Way Pipes



Pipes with Fork (I)



```
main() {  
    int  childpid, pipe1[2], pipe2[2];  
    if (pipe(pipe1) < 0 || pipe(pipe2) < 0) err_sys("can't create pipes");  
    if ( (childpid = fork()) < 0) {  
        err_sys("can't fork");  
    }
```


Pipes with Fork (II)

```
    } else if (childpid > 0) {                                /* parent */
        close(pipe1[0]); close(pipe2[1]);
        client(pipe2[0], pipe1[1]);
        while (wait((int *) 0) != childpid)                  /* wait for child */
            ;
        close(pipe1[1]); close(pipe2[0]);
        exit(0);
    } else {                                                  /* child */
        close(pipe1[1]); close(pipe2[0]);
        server(pipe1[0], pipe2[1]);
        close(pipe1[0]); close(pipe2[1]);
        exit(0);
    }
}
```

Client

```
client(int readfd, int writefd)
{
    char buff[MAXBUFF];
    int n;

    // Read the filename from standard input, write it to the IPC descriptor.
    if (fgets(buff, MAXBUFF, stdin) == NULL) err_sys( "client: filename read error" );

    n = strlen(buff);
    if (buff[n-1] == '\n' ) n--; /* ignore newline from fgets() */
    if (write(writefd, buff, n) != n) err_sys('client: filename write error');

    // Read the data from the IPC descriptor and write to standard output.
    while ( (n = read(readfd, buff, MAXBUFF)) > 0)
        if (write(1 /* stdout*/, buff, n) != n) err_sys('client: data write error');
    if (n < 0) err_sys('client: data read error');
}
```

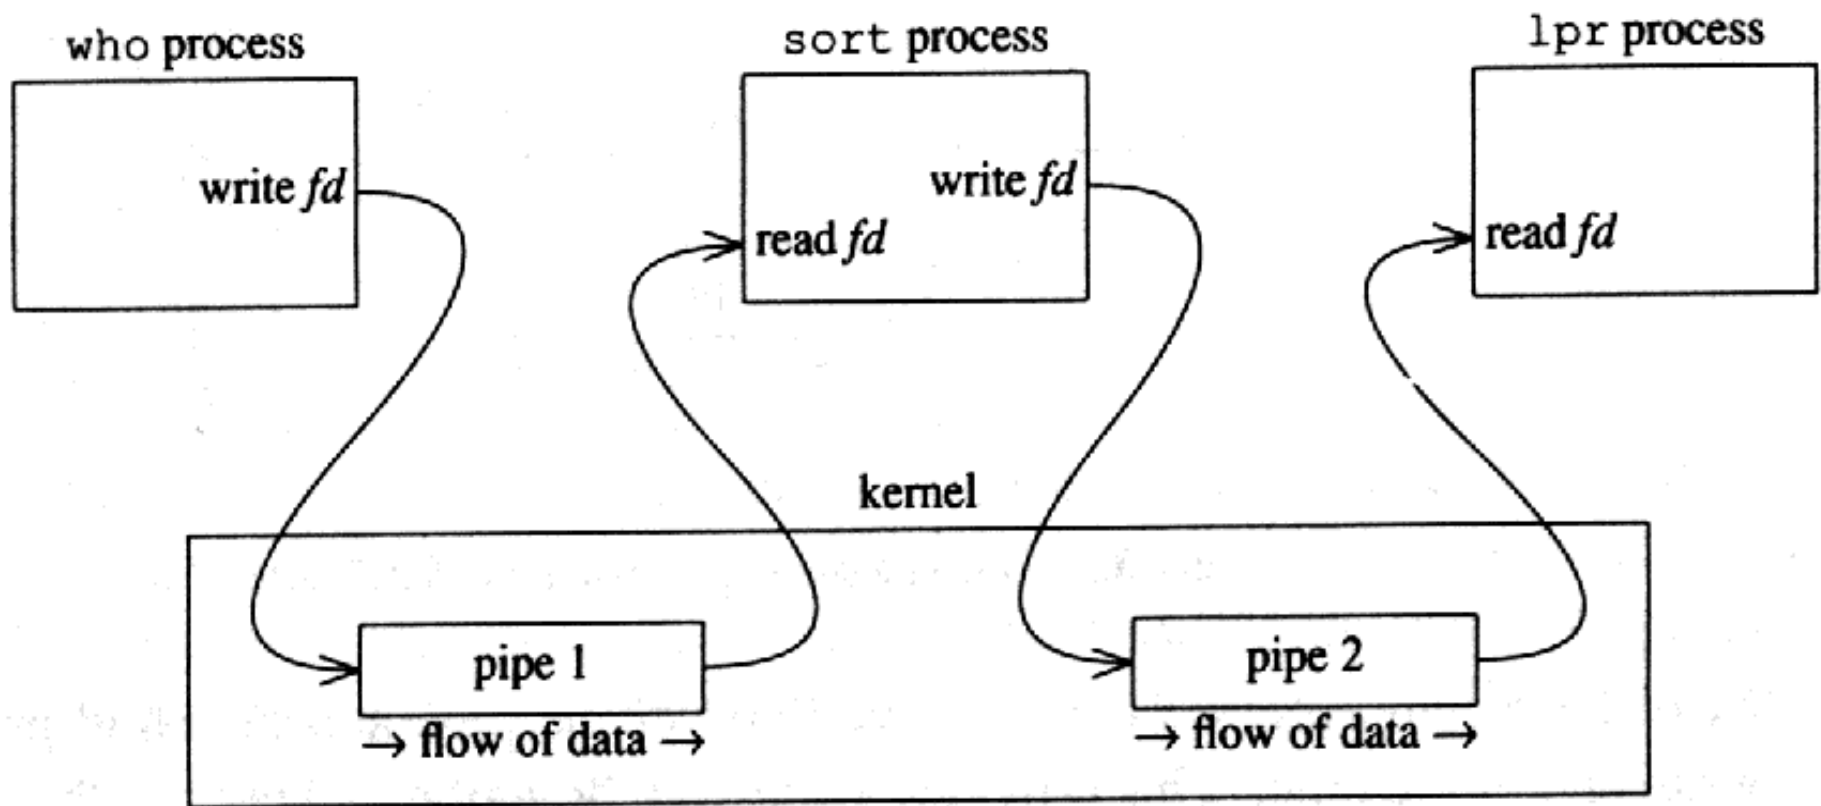
Server

```
server(int readfd, int writefd)
{
    char    buff[MAXBUFF], errmsg[256], *sys_err_str();
    int     n, fd;

    // Read the filename from the IPC descriptor.
    if ( (n = read(readfd, buff, MAXBUFF)) <= 0) err_sys("server: filename read error");
    buff[n] = '\0';                          /* null terminate filename */

    if ( (fd = open(buff, 0)) < 0) {
        // Error. Format an error message and send it back to the client.
        sprintf(errmsg, ": can't open, %s\n", sys_err_str());
        strcat(buff, errmsg);
        n = strlen(buff);
        if (write(writefd, buff, n) != n) err_sys("server: errmsg write error");
    } else {
        // Read the data from the file and write to the IPC descriptor.
        while ( (n = read(fd, buff, MAXBUFF)) > 0)
            if (write(writefd, buff, n) != n) err_sys("server: data write error");
        if (n < 0) err_sys("server: read error");
    }
}
```

Unix Shell Pipes



SIGPIPE

What if none of processes read or write a pipe?

- When trying to read from a pipe that has no write end,
 - the **read()** returns 0.
- When you try to write to a pipe that has no read end,
 - a **SIGPIPE** signal is generated.
 - If the signal isn't handled, the program exits silently.
 - A nice way for the example

```
UNIX> cat exec1.c | head -5 | tail -1
```

 - ▶ **head** exits when receiving 5 lines,
 - ▶ **tail** will have **read()** return 0, and will exit, and
 - ▶ **cat** will try to write to an empty pipe, and thus will generate **SIGPIPE** and exit.

See <http://www.cs.utk.edu/~plank/plank/classes/cs360/360/notes/Pipe/lecture.html>

FIFOs

Stands for first-in-first-out.(System V only)

Main problem of pipes:

- can only be used between processes that have a parent process in common.
- Server
 `mknod(pathname, mode, dev)`
- Client
 `open(pathname, flag)`

Example: FIFO Server

```
#define FIFO1    "/tmp/fifo.1"
#define FIFO2    "/tmp/fifo.2"
main() {
    int    readfd, writefd;
    // Create the FIFOs, then open them - one for reading and one for writing.
    if ( (mknod(FIFO1, S_IFIFO | PERMS, 0) < 0) && (errno != EEXIST))
        err_sys("can't create fifo: %s", FIFO1);
    if ( (mknod(FIFO2, S_IFIFO | PERMS, 0) < 0) && (errno != EEXIST)) {
        unlink(FIFO1);
        err_sys("can't create fifo: %s", FIFO2);
    }
    if ( (readfd = open(FIFO1, 0)) < 0)
        err_sys("server: can't open read fifo: %s", FIFO1);
    if ( (writefd = open(FIFO2, 1)) < 0)
        err_sys("server: can't open write fifo: %s", FIFO2);
    server(readfd, writefd);
    close(readfd);
    close(writefd);
}
```

Example: FIFO Client

```
main() {  
    int    readfd, writefd;  
    // Open the FIFOs. We assume the server has already created them.  
    if ( (writefd = open(FIFO1, 1)) < 0)  
        err_sys("client: can't open write fifo: %s", FIFO1);  
    if ( (readfd = open(FIFO2, 0)) < 0)  
        err_sys("client: can't open read fifo: %s", FIFO2);  
    client(readfd, writefd);  
    close(readfd);  
    close(writefd);  
    // Delete the FIFOs, now that we're finished.  
    if (unlink(FIFO1) < 0) err_sys("client: can't unlink %s", FIFO1);  
    if (unlink(FIFO2) < 0) err_sys("client: can't unlink %s", FIFO2);  
}
```


FIFOs over Two Hosts

What if Client (Host) A makes file mounting on Server B?

- Create a FIFO on B.
- Both processes on A can communicate via it correctly.
- One on A and the other one B cannot communicate.
- For different Clients:

Processes A&B run on	Does the FIFO work?
same client	Yes
different client	No

- References:

- <http://hissa.nist.gov/rbac/titleissues/node24.html>
- <https://mail.rtai.org/pipermail/rtai/2006-July/015566.html>

Message Queues

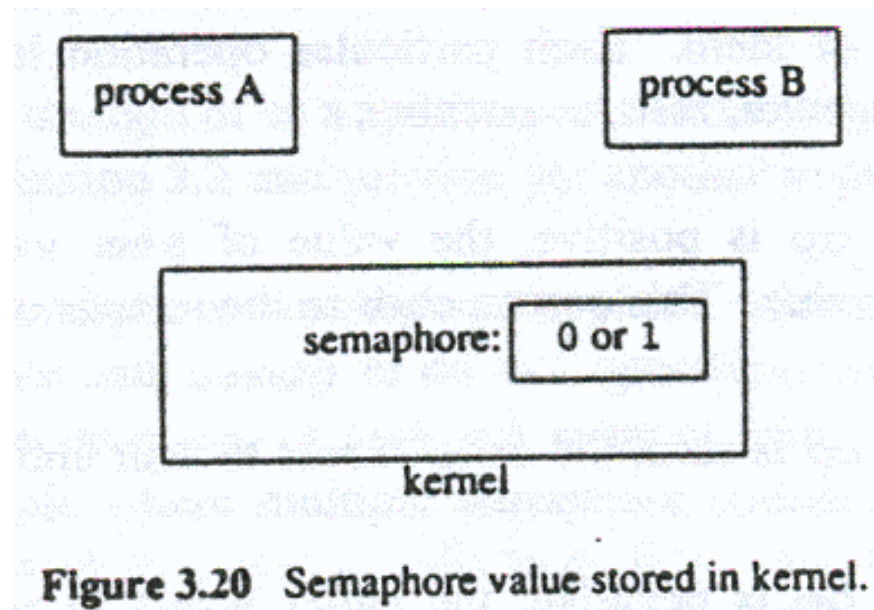
- Message queues: (supported by System V)

Message v.s. Stream:

- Message: has boundary.
- Stream: no boundary for two messages.

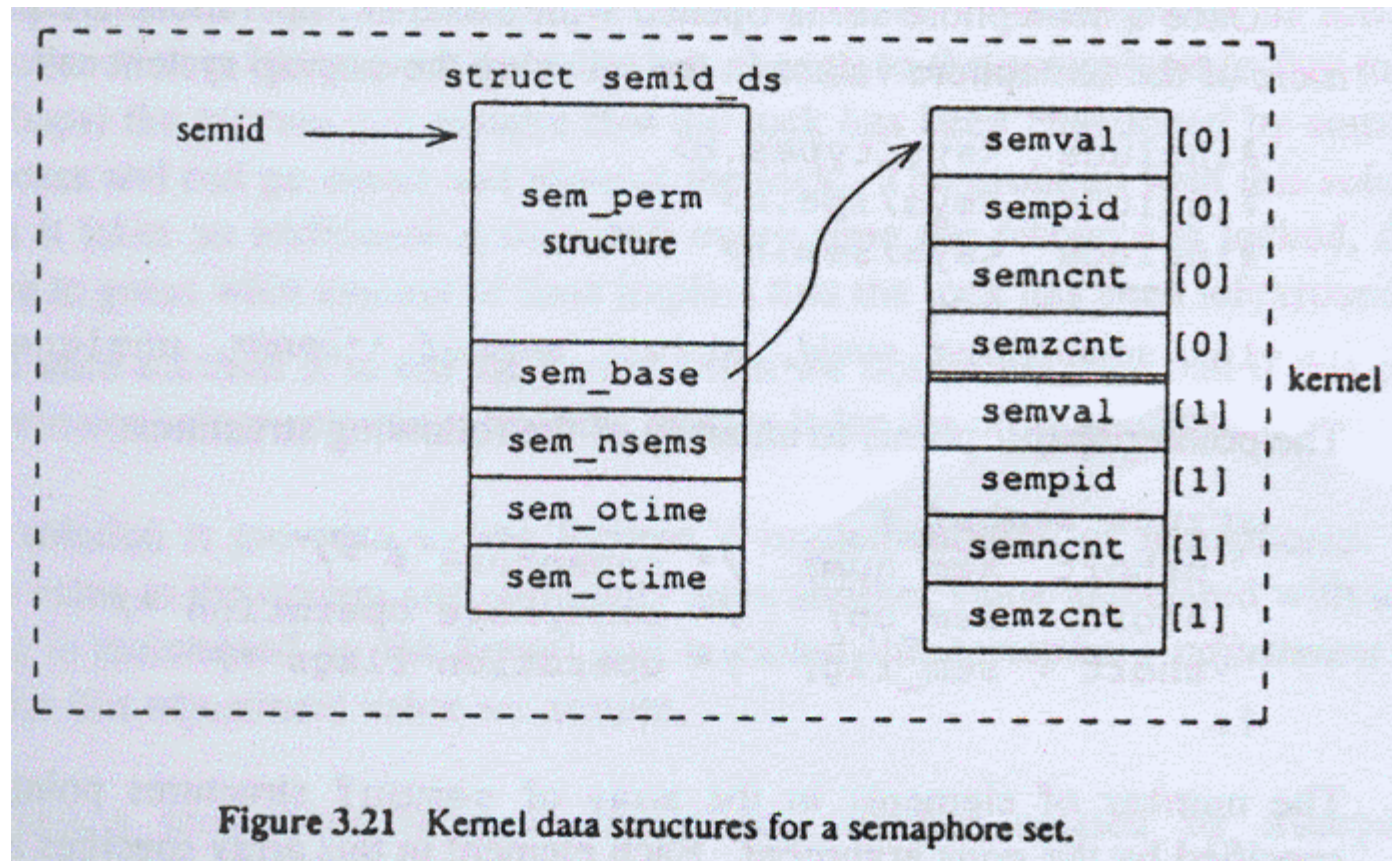
Semaphores

- This does not really send data as read/write/pipe/fifo, simply does synchronous operations to protect shared memory.
- Semaphores are stored in kernel.



Kernel Data Structures for Semaphores

- Usually, a set of semaphores (in arrays), not just one.



Allocating Semaphores

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
```

```
int semget(key_t key, int nsems, int semflag);
```

Numeric	Symbolic	Description
0400	SEM_R	Read by owner
0200	SEM_A	Alter by owner
0040	SEM_R >> 3	Read by group
0020	SEM_A >> 3	Alter by group
0004	SEM_R >> 6	Read by world
0002	SEM_A >> 6	Alter by world
	IPC_CREAT	(See Section 3.8)
	IPC_EXCL	(See Section 3.8)

semop

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
```

```
int semop(int semid, struct sembuf * opsptr, unsigned int nops);
```

pointer *opsptr* points to an array of the following structures:

```
struct sembuf {
    ushort  sem_num;    /* semaphore # */    Note: more like ID.
    short   sem_op;     /* semaphore operation */
    short   sem_flg;    /* operation flags */
};
```

Semaphore Operations

- If `sem_op` is positive, the value of `sem_val` is added.
- If `sem_op` is zero, the caller wants to wait until the semaphore becomes zero.
- If `sem_op` is negative, the caller wants to wait until the semaphore becomes greater than or equal to the absolute value of `sem_op`. Then, add the value into `sem_val`.
(Just like subtract)

Returning value:

- 0, if ok.
- -1, for an error.

File Locking with Semaphore

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

#define SEMKEY 123456L /* key value for semget() */
#define PERMS 0666

static struct sembuf op_lock[2] = {
    0, 0, 0,          /* wait for sem#0 to become 0 */
    0, 1, 0,          /* then increment sem#0 by 1 */
};

static struct sembuf op_unlock[1] = {
    0, -1, IPC_NOWAIT /* decrement sem#0 by 1 (sets it to 0) */
};

int semid = -1; /* semaphore id */
```


File Locking with Semaphore (cont.)

```
my_lock(int fd)
{
    if (semid < 0) {
        if ( (semid = semget(SEMKEY, 1, IPC_CREAT | PERMS)) < 0)
            err_sys("semget error");
    }
    if (semop(semid, &op_lock[0], 2) < 0)
        err_sys("semop lock error");
}

my_unlock(fd)
int fd;
{
    if (semop(semid, &op_unlock[0], 1) < 0)
        err_sys("semop unlock error");
}
```

Problems

- If a process aborts while locking,
 - the semaphore value remains one!!!

Possible solutions:

- Catch all signals and use signal handlers to unlock it.
 - But, SIGKILL cannot be caught.
- Let the first `sem_op` not wait, get the `sem_ctime` to check if timeout happens.
 - Too sophisticated and does not solve the problem really.
- Let the kernel undo it, while the process aborts.
 - Use “SEM_UNDO” flag to tell kernel to undo.

File Locking with Semaphore Undo

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

#define SEMKEY 123456L /* key value for semget() */
#define PERMS 0666

static struct sembuf op_lock[2] = {
    0, 0, 0,          /* wait for sem#0 to become 0 */
    0, 1, SEM_UNDO    /* then increment sem#0 by 1 */
};

static struct sembuf op_unlock[1] = {
    0, -1, (IPC_NOWAIT | SEM_UNDO)
    /* decrement sem#0 by 1 (sets it to 0) */
};
```

Problems Still

- The semaphore is never removed.
 - Though we can use `semctl()` to remove it, a process that aborts may not have chance to call this.
 - The `my_lock()` code between `semget()` and `semop()` is not atomic.

semctl

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

int semctl(int semid, int semnum, int cmd, union semun arg);

union semun {
    int          val;      /* used for SETVAL only */
    struct semid_ds *buff; /* used for IPC_STAT and IPC_SET */
    ushort       *array;   /* used for IPC_GETALL & IPC_SETALL */
} arg;
```

- Remove a lock.
 `semctl(semid, 0, IPC_RMID, (struct semun *) 0)`
- Get and Set semaphore values.
 `semval = semctl(id, 1, GETVAL, 0)`
 `semctl(id, 2, SETVAL, semctl_arg)`

A Robust Semaphore

- Use 3 semaphore values
 1. The real semaphore value
 2. The counter of the number of processes using this semaphore.
 3. A lock variable for the semaphore. (Used for a critical section in code.)

- Provide a simpler and easier to understand interface: 7 routines

<code>id = sem_create(key, initval);</code>	<code># create with initial value or open</code>
<code>id = sem_open(key);</code>	<code># open (must already exist)</code>
<code>sem_wait(id);</code>	<code># wait = P = down by 1</code>
<code>sem_signal(id);</code>	<code># signal = V = up by 1</code>
<code>sem_op(id, amount);</code>	<code># wait if (amount < 0)</code>
	<code># signal if (amount > 0)</code>
<code>sem_close(id);</code>	<code># close</code>
<code>sem_rm(id);</code>	<code># remove (delete)</code>

sem_create() — (1)

```
int sem_create(key_t key, int initval)
{
    register int          id, semval;
    if (key == IPC_PRIVATE) return(-1); /* not intended for private semaphores */
    else if (key == (key_t) -1) return(-1); /* probably an ftok() error by caller */
```

again:

```
    if ( (id = semget(key, 3, 0666 | IPC_CREAT)) < 0) return(-1);
    if (semop(id, &op_lock[0], 2) < 0) {
        if (errno == EINVAL) goto again;
        err_sys("can't lock");
```


```
    }
```

```
static struct sembuf op_lock[2] = {
    2, 0, 0, /* wait for [2] (lock) to equal 0 */
    2, 1, SEM_UNDO /* then increment [2] to 1 - this locks it */
                /* UNDO to release the lock if processes exits
                before explicitly unlocking */
```

sem_create() — (2)

```
if ( (semval = semctl(id, 1, GETVAL, 0)) < 0) err_sys("can't GETVAL");
if (semval == 0) {
    semctl_arg.val = initval;
    if (semctl(id, 0, SETVAL, semctl_arg) < 0)
        err_sys("can't SETVAL[0]");
    semctl_arg.val = BIGCOUNT;
    if (semctl(id, 1, SETVAL, semctl_arg) < 0)
        err_sys("can't SETVAL[1]");
}
if (semop(id, &op_endcreate[0], 2) < 0)
    err_sys("can't end create");

return(id);
}
```



```
static struct sembuf op_endcreate[2] = {
    1, -1, SEM_UNDO, /* decrement [1] (proc counter) with undo on exit */
                                /* UNDO to adjust proc counter if process exits
                                before explicitly calling sem_close() */
    2, -1, SEM_UNDO    /* decrement [2] (lock) back to 0 -> unlock */
};
```


sem_rm()

```
sem_rm(int id)
{
    if (semctl(id, 0, IPC_RMID, 0) < 0)
        err_sys("can't IPC_RMID");
}
```

sem_open()


```
int sem_open(key_t key)
{
    register int id;
    if (key == IPC_PRIVATE) return(-1);
    else if (key == (key_t) -1) return(-1);

    if ( (id = semget(key, 3, 0)) < 0) return(-1);    /* doesn't exist, or tables full */

    // Decrement the process counter. We don't need a lock to do this.
    if (semop(id, &op_open[0], 1) < 0) err_sys("can't open");

    return(id);
}

static struct sembuf op_open[1] = {
    1, -1, SEM_UNDO    /* decrement [1] (proc counter) with undo on exit */
};
```



sem_close()

```
sem_close(int id)
```

```
{
```

```
    register int semval;
```

```
    // The following semop() first gets a lock on the semaphore,
```

```
    // then increments [1] - the process counter.
```

```
    if (semop(id, &op_close[0], 3) < 0) err_sys("can't semop");
```

```
    // if this is the last reference to the semaphore, remove this.
```

```
    if ( (semval = semctl(id, 1, GETVAL, 0)) < 0) err_sys("can't GETVAL");
```

```
    if (semval > BIGCOUNT) err_dump("sem[1] > BIGCOUNT");
```

```
    else if (semval == BIGCOUNT) sem_rm(id);
```

```
    else
```

```
        if (semop(id, &op_unlock[0], 1) < 0) err_sys("can't unlock"); /* unlock */
```

```
}
```

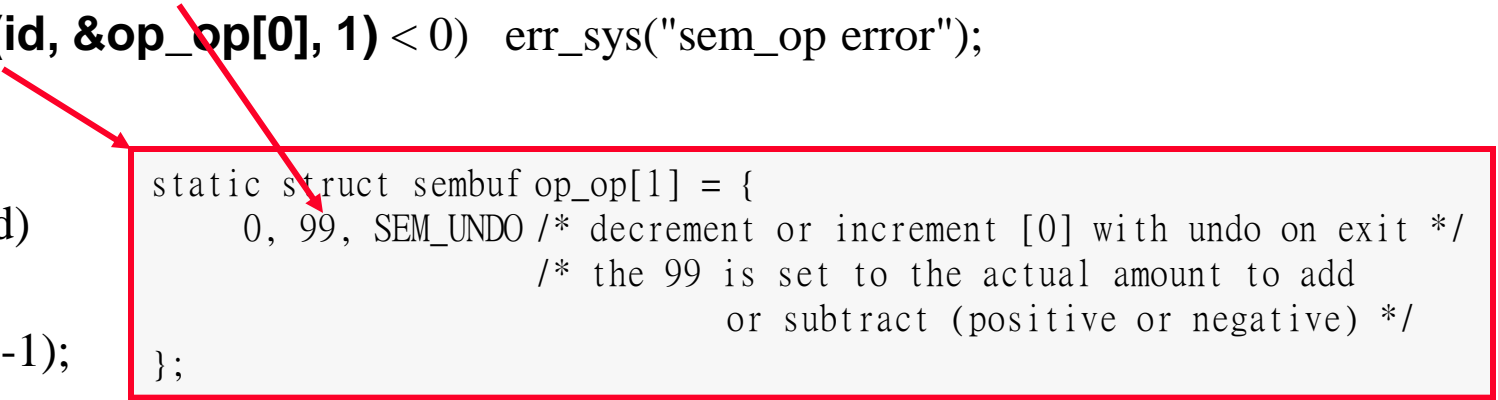
```
static struct sembuf op_close[3] = {
    2, 0, 0, /* wait for [2] (lock) to equal 0 */
    2, 1, SEM_UNDO, /* then increment [2] to 1 - this locks it */
    1, 1, SEM_UNDO /* then increment [1] (proc counter) */
};
```

```
static struct sembuf op_unlock[1] = {
    2, -1, SEM_UNDO /* decrement [2] (lock) back to 0 */
};
```

P/V Operations

```
sem_op(int id, int value)
{
    if ( (op_op[0].sem_op = value) == 0) err_sys("can't have value == 0");
    if (semop(id, &op_op[0], 1) < 0) err_sys("sem_op error");
}
```

```
sem_wait(int id)
{
    sem_op(id, -1);
}
```



```
static struct sembuf op_op[1] = {
    0, 99, SEM_UNDO /* decrement or increment [0] with undo on exit */
    /* the 99 is set to the actual amount to add
       or subtract (positive or negative) */
};
```

```
sem_signal(int id)
{
    sem_op(id, 1);
}
```

Locking with Semaphores

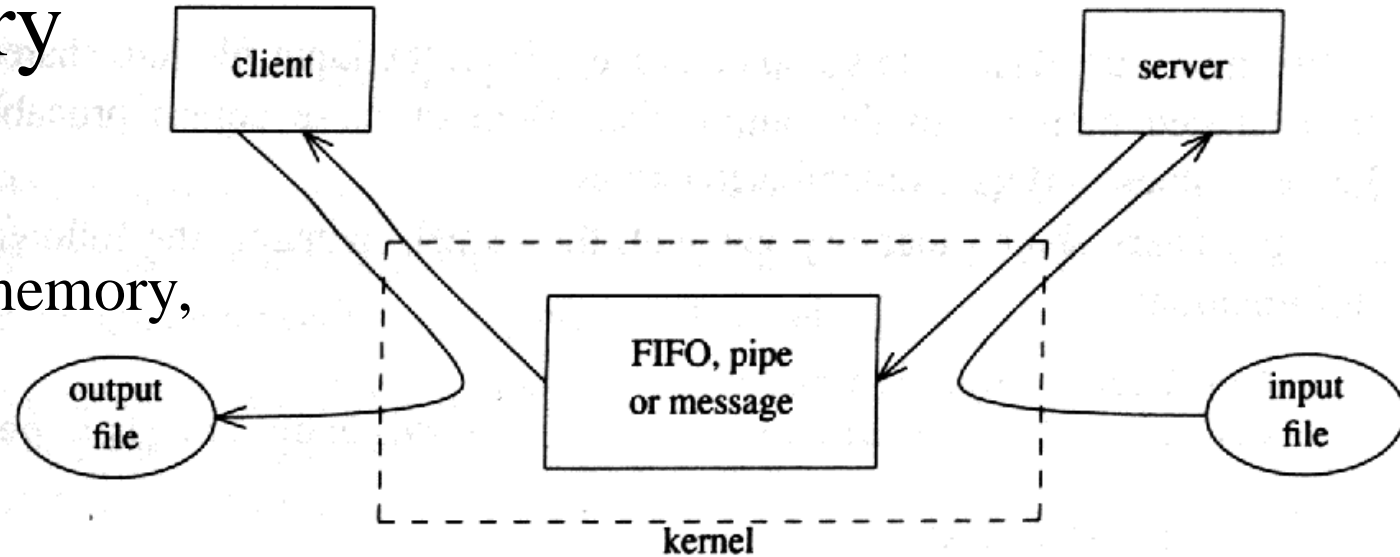
```
#define      SEMKEY      123456L /* key value for sem_create() */
int semid = -1; /* semaphore id */

my_lock(int fd)
{
    if (semid < 0) {
        if ( (semid = sem_create(SEMKEY, 1)) < 0) err_sys("sem_create error");
    }
    sem_wait(semid);
}


my_unlock(int fd)
{
    sem_signal(semid);
}
```

Shared Memory

- Without shared memory, too many copies.



- Our next example:



System Call – shmget

- Create a shared memory segment or access an existing one.

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

int shmget(key_t key, int size, int shmflag);
```

- Return *shmid*, or -1 for error.
- *size*: the size of the segment
- *shmflag*: the flags listed right.

Numeric	Symbolic	Description
0400	SHM_R	Read by owner
0200	SHM_W	Write by owner
0040	SHM_R >> 3	Read by group
0020	SHM_W >> 3	Write by group
0004	SHM_R >> 6	Read by world
0002	SHM_W >> 6	Write by world
	IPC_CREAT	(See Section 3.8)
	IPC_EXCL	(See Section 3.8)

System Call – shmat

- Attach the shared memory segment.

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

char *shmat(int shmid, char *shmaddr, int shmflag);
```

- *shmaddr*:

- 0: the system selects the address for the caller.
- Nonzero:
 - ▶ SHM_RND value is not specified:
 - attached at the specified address, *shmaddr*.
 - ▶ SHM_RND value is specified:
 - attached at the specified address, *shmaddr*, but rounded down by SHMLBA (in *shmflag*).

- *shmflag*:

- SHM_RDONLY: “read-only” access.
- SHMLBA (See above).

System Call – shmdt

- Detach the shared memory segment.

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

int shmdt(char *shmaddr);
```

- This does not really delete the shared memory segment.

System Call – shmctl

- To remove a shared memory, use this with cmd “IPC_RMID”.

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

int shmctl(int shmid, int cmd, struct shm_id_ds *buf);
```

Header File – shm.h

```
#define      SHMKEY          ((key_t) 7890) /* base value for shmem key */
#define      SEMKEY1        ((key_t) 7891) /* client semaphore key */
#define      SEMKEY2        ((key_t) 7892) /* server semaphore key */
#define      PERMS    0666
int shmid, clisem, servsem;          /* shared memory and semaphore IDs */

#define      MAXMESGDATA (4096-16)
#define      MSGHDRSIZE  (sizeof(Mesg) - MAXMESGDATA)
typedef struct {
    int  mesg_len;          /* #bytes in mesg_data, can be 0 or > 0 */
    long mesg_type;         /* message type, must be > 0 */
    char mesg_data[MAXMESGDATA];
} Mesg;
Mesg      *mesgptr;        /* ptr to message structure, which is
                           in the shared memory segment */
```

Client – main()

```
main() {  
    // Get the shared memory segment and attach it.  
    // The server must have already created it.  
    if ( (shmid = shmget(SHMKEY, sizeof(Mesg), 0)) < 0) err_sys("...");  
    if ( (mesgptr = (Mesg *) shmat(shmid, (char *) 0, 0)) == -1) err_sys("...");  
    // Open the two semaphores. The server must have created them already.  
    if ( (clisem = sem_open(SEMKEY1)) < 0) err_sys("...");  
    if ( (servsem = sem_open(SEMKEY2)) < 0) err_sys("...");  
    client();  
    // Detach and remove the shared memory segment and close the semaphores.  
    if (shmdt(mesgptr) < 0) err_sys("...");  
    if (shmctl(shmid, IPC_RMID, (struct shmid_ds *) 0) < 0) err_sys("...");  
    sem_close(clisem);           /* will remove the semaphore */  
    sem_close(servsem);        /* will remove the semaphore */  
    exit(0);  
}
```

Server – main()

```
main() {  
    // Create the shared memory segment, if required, then attach it.  
    if ((shmid=shmget(SHMKEY, sizeof(Mesg), PERMS|IPC_CREAT))<0)  
        err_sys("server: can't get shared memory");  
    if ( (mesgptr = (Mesg *) shmat(shmid, (char *) 0, 0)) == -1) err_sys("...");  
    // Create two semaphores. The client semaphore starts out at 1  
    // since the client process starts things going.  
    if ( (clisem = sem_create(SEMKEY1, 1)) < 0) err_sys("...");  
    if ( (servsem = sem_create(SEMKEY2, 0)) < 0) err_sys("...");  
    server();  
    // Detach the shared memory segment and close the semaphores.  
    // The client is the last one to use the shared memory, so it'll remove it at last.  
    if (shmdt(mesgptr) < 0) err_sys("server: can't detach shared memory");  
    sem_close(clisem);  
    sem_close(servsem);
```

Client – client()

```
client() {  
    // Read the filename from standard input, write it to shared memory.  
    sem_wait(clisem);                /* get control of shared memory */  
    if (fgets(msgptr->msg_data, MAXMSGDATA, stdin) == 0) err_sys("...");  
    n = strlen(msgptr->msg_data);  
    if (msgptr->msg_data[n-1] == '\n') n--; /* ignore newline from fgets() */  
    msgptr->msg_len = n;  
    sem_signal(servsem);            /* wake up server */  
    // Wait for the server to place something in shared memory.  
    sem_wait(clisem);                /* wait for server to process */  
    while( (n = msgptr->msg_len) > 0) {  
        if (write(1, msgptr->msg_data, n) != n) err_sys("data write error");  
        sem_signal(servsem);        /* wake up server */  
        sem_wait(clisem);            /* wait for server to process */  
    }  
    if (n < 0) err_sys("data read error");  
}
```

Server – server()

```
server() {  
    int          n, filefd;  
    char         errmsg[256], *sys_err_str();  
  
    // Wait for the client to write the filename into shared memory.  
    sem_wait(servsem);           /* we'll wait here for client to start things */  
    msgptr->msg_data[msgptr->msg_len] = '\0';  
  
    if ( (filefd = open(msgptr->msg_data, 0)) < 0) {  
        // Error. Format an error message and send it back to the client.  
        sprintf(errmsg, ": can't open, %s\n", sys_err_str());  
        strcat(msgptr->msg_data, errmsg);  
        msgptr->msg_len = strlen(msgptr->msg_data);  
        sem_signal(clisem);           /* send to client */  
        sem_wait(servsem);           /* wait for client to process */  
    }
```

Server – server() (cont.)

```
} else {  
    // Read the data from the file right into shared memory.  
    while ( (n = read(filefd, msgptr->msg_data, MAXMSGDATA-1)) >  
0) {  
        msgptr->msg_len = n;  
        sem_signal(clisem);           /* send to client */  
        sem_wait(servsem);           /* wait for client to process */  
    }  
    close(filefd);  
    if (n < 0) err_sys("server: read error");  
}  
// Send a message with a length of 0 to signify the end.  
msgptr->msg_len = 0;  
sem_signal(clisem);  
}
```